

Subline

Subline draws per-token sparklines directly beneath the text they describe. Each segment is aligned to its token, so you can see at a glance which tokens carry the most information – without losing the context of the surrounding text.

The visualization adapts to your system light/dark preference automatically.

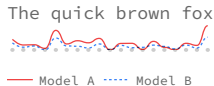
Single series

The simplest case: one value per token, plotted as a sparkline beneath each character. Here we show a rough hand-crafted surprise signal for the sentence "The quick brown fox." Spaces and common letters score low; the rarer "q" and "x" score high.



Multiple series

Multiple series overlay on the same sparkline, letting you compare two models side by side, or pair a metric with its complement. A dashed line keeps them visually distinct.



Wrapping long text

When the text exceeds `chars_per_line`, Subline wraps it automatically. Each line gets its own sparkline, and the token alignment is preserved across the break.

